

Transformer models have become the default architecture for large-scale language processing. Residual connections and layer normalization stabilize training at depth. Paragraph 1.1 of the scanned report.

Attention computes a weighted combination of values. Scaled dot-product attention remains the most common building block in encoder and decoder stacks. Paragraph 1.2 of the scanned report.

Positional encodings inject order information that self-attention cannot recover from the input tokens alone. Relative variants later improved length generalization. Paragraph 1.3 of the scanned report.

Transformer models have become the default architecture for large-scale language processing. Residual connections and layer normalization stabilize training at depth. Paragraph 2.1 of the scanned report.

Attention computes a weighted combination of values. Scaled dot-product attention remains the most common building block in encoder and decoder stacks. Paragraph 2.2 of the scanned report.

Positional encodings inject order information that self-attention cannot recover from the input tokens alone. Relative variants later improved length generalization. Paragraph 2.3 of the scanned report.

Transformer models have become the default architecture for large-scale language processing. Residual connections and layer normalization stabilize training at depth. Paragraph 3.1 of the scanned report.

Attention computes a weighted combination of values. Scaled dot-product attention remains the most common building block in encoder and decoder stacks. Paragraph 3.2 of the scanned report.

Positional encodings inject order information that self-attention cannot recover from the input tokens alone. Relative variants later improved length generalization. Paragraph 3.3 of the scanned report.