

Body of page 1. Transformer models have become the default architecture for large-scale language processing. Residual connections and layer normalization stabilize training at depth.

Body of page 2. Transformer models have become the default architecture for large-scale language processing. Residual connections and layer normalization stabilize training at depth.